

UML501

Machine Learning Project

Submitted By:

Arnav Srivastava – 102303120

Ayush Kaushik – 102303104



Submitted To:

Ms. Reeta

Computer Science Engineering Department

December 2025

INDEX

Content	Page No.
1. Introduction	3-5
2. Literature Survey	6-8
3. Research Gap	9-10
4. Methodology	11-13
5. Results and Discussion	14-19
6. Conclusion	20-21
7. Future Work	22-23
8. References	24

INTRODUCTION

Social media platforms like Twitter (now X) and Facebook have revolutionized communication, allowing users to share opinions instantly. However, this freedom of expression has led to a significant rise in hate speech and offensive language. Hate speech refers to abusive or threatening speech or writing that expresses prejudice against a particular group, especially based on race, religion, or sexual orientation.

Hate speech has become one of the most critical challenges in today's digital communication landscape. With the rapid growth of social media platforms, people can express their opinions instantly and anonymously, resulting in an increase in harmful and offensive content. Such content is not only offensive but can also incite violence, discrimination, or hostility within society.

The sheer volume of user-generated content makes manual moderation impossible. Consequently, automated systems are required to filter toxic content to maintain a healthy online environment. Traditional keyword-based filtering systems often fail to understand context (e.g., distinguishing between offensive language used colloquially versus hateful intent).

This project focuses on developing a **Hate Speech Detection System** using Natural Language Processing (NLP) and Deep Learning techniques. The dataset contains labelled text categorized into hate speech, offensive language, or normal speech. The goal is to preprocess text, extract features using TF-IDF/word embeddings, and train ML/DL models to classify messages accurately.

The system is implemented using a **Bi-Directional LSTM (Long Short-Term Memory)** neural network, along with comparative analysis using Machine Learning models like **XGBoost** and **Logistic Regression**. Developing robust classification models helps create safer digital environments, assists content moderators, and promotes healthier online communication.

The implemented system demonstrates how artificial intelligence can automate the identification of hate speech and offensive language from textual data without requiring labour-intensive manual moderation. Rather than relying on static keyword filters or human intervention, the solution provides a Deep Learning-based approach. This makes the proposed solution scalable and highly effective for maintaining a safe online environment across vast social media platforms.

The goal is to classify tweets into three specific categories:

1. **Hate Speech** (Class 0)
2. **Offensive Language** (Class 1)
3. **Neither** (Class 2)

By leveraging techniques like TF-IDF vectorization and Word Embeddings, this project demonstrates how AI can automate content moderation with high accuracy.

LITERATURE SURVEY

The proliferation of hate speech and offensive content across social media platforms (such as Twitter, Facebook, and YouTube) has made automated detection a critical research area in Natural Language Processing (NLP). This survey reviews key advancements in deep learning models and data handling techniques employed to address this challenge, particularly focusing on the role of Recurrent Neural Networks (RNNs) like Bi LSTM and the state-of-the-art Transformer architectures.

1. The Necessity of Automated Detection and Application Context

The sheer volume and speed of content generated on social media necessitate automated tools for real-time moderation, a complex task due to the ambiguous, context-dependent, and evolving nature of hateful language. The primary goal is to classify textual content into categories such as:

- **Hate Speech:** Content explicitly targeting individuals or groups based on protected characteristics (e.g., race, religion, gender).
- **Offensive Language/Abusive Language:** Content containing profanity or general insults, but not necessarily targeting a protected group.

- **Neither/Normal:** Non-hateful, non-offensive text.

2. Evolution from Traditional Machine Learning (ML)

Early hate speech detection research relied heavily on feature engineering and traditional machine learning algorithms

3. Deep Learning Architectures for Contextual Understanding

Deep Learning (DL) models are preferred as they can automatically learn complex, high-level feature representations from raw text, effectively addressing the limitations of traditional, hand-crafted features.

A. Recurrent Neural Networks (RNNs) and Bi-LSTM

Bidirectional Long Short-Term Memory (Bi-LSTM) networks have been widely adopted for sequence classification tasks like hate speech detection.

B. Transformer-based Architectures (BERT)

The current state-of-the-art in text classification is dominated by pre-trained language models based on the **Transformer** architecture, such as **BERT (Bidirectional Encoder Representations from Transformers)**.

4. Critical Challenge: Handling Imbalanced Datasets

Hate speech datasets are intrinsically imbalanced, as the majority of content is non-hateful (the majority class), while the critical **Hate**

Speech class is a severe minority. This imbalance causes classifiers to become biased towards the majority class, leading to high overall accuracy but poor performance (low recall/F1-score) on the minority "Hate Speech" class.

Data Balancing Techniques

To mitigate this issue, researchers implement data-level and algorithmic-level solutions:

- **Sampling-based Methods:**
 - **Oversampling:** Increasing the size of the minority class (Hate Speech) using methods like Random Oversampling (ROS) or **SMOTE (Synthetic Minority Oversampling Technique)**. This balances the training data distribution.
 - **Undersampling:** Reducing the size of the majority class (Non-Hate Speech).
- **Cost-Sensitive Learning:** Adjusting the loss function during model training to penalize misclassification of the minority class more heavily than the majority class.

RESEARCH GAP

1.1 Research Gaps

Despite extensive research in text detection systems and voice detection, several gaps remain in the domain of **hate speech detection and priority handling**.

1. Contextual Ambiguity

Existing systems often struggle to differentiate between "Offensive Language" (profanity used casually) and "Hate Speech" (targeted attacks). Many models classify any profanity as hate speech, leading to false positives.

2. Data Imbalance

Real-world datasets are heavily skewed. In the dataset used for this project, "Offensive Language" instances vastly outnumbered "Hate Speech." Standard training on such data leads to models that ignore the minority class (Hate Speech).

3. Model Generalization

While large language models (like BERT) exist, they are computationally expensive. There is a need for lightweight yet accurate Deep Learning models (like BiLSTMs) that can run efficiently without requiring massive hardware resources.

4. Limited Dataset Diversity

Public datasets for hate speech detection are limited in quantity and diversity. Existing models often suffer from class imbalance (for example, fewer hate speech tweets than offensive or neutral tweets), which affects detection accuracy and model generalization.

METHODOLOGY

4. Methodology

This project follows a computer text-based deep learning approach for detecting foul and hateful language from tweets. The workflow includes dataset preparation, model training, evaluation, and application deployment.

4.1 Dataset Description

Dataset Preparation The dataset consists of tweets labeled into three classes:

- Hate Speech (0)
- Offensive (1)
- Neither (2)
-
- **Initial Analysis:** The raw data was highly imbalanced, with Class 1 (Offensive) dominating.
- **Balancing:** We applied a resampling strategy. Class 0 and Class 2 were oversampled, and Class 1 was downsampled to create a balanced dataset of **4,163 samples per class** (Total: ~12,500 samples).

4.2 Data Preprocessing

To make the text suitable for the model, the following NLP techniques were applied:

1. **Cleaning:** Conversion to lowercase, removal of punctuation and special characters.
2. **Stopword Removal:** Removing common words (e.g., "the", "is") that do not add semantic value.

3. **Lemmatization:** Reducing words to their base root form (e.g., "running" -> "run") using WordNetLemmatizer.
4. **Tokenization:** Converting text into numerical sequences. A vocabulary size of 10,000 words was used.
5. **Padding:** All sequences were padded to a fixed length of 100 to ensure a uniform input shape.

The dataset was split into:

- 80% training
- 20% testing

4.3 Model Training

- **Bi-LSTM** was chosen as the **core classification model** because it is a highly effective deep learning architecture for processing sequential data like text.
- **Logistic Regression (LR)** was selected as the **linear baseline model** for comparison.
- **XGBoost (Extreme Gradient Boosting)** was selected as the **high-performance machine learning benchmark**.

4.4 Model Evaluation

Performance was measured using:

- Precision
- Recall
- Confusion matrix
- Accuracy
- Graphs

4.5 Application Implementation

The system was implemented using:

- Python
- Google Colab
- Hate-speech detection dataset

4.6 Advantages

- Fully automated detection
- Fostering a Safer and More Inclusive Online Environment
- Scalability and Efficiency
- Support for Platform Credibility and Legal Compliance
- Advancing NLP and Deep Learning Research

RESULTS AND DISCUSSION

This section presents the experimental results obtained after training and testing the Hate Speech Detection model. The performance of the system is evaluated based on standard text classification metrics and observed behavior during testing with text (tweets).


5.1 Detection Performance

The trained models demonstrated strong performance in detecting hate speech from real-world tweets.

Performance Metrics

The model was evaluated using scikit-learn evaluation metrics, which provide the following metrics:

- Precision
- Accuracy
- Recall
- Confusion Matrix



BiLSTM Model :

```
... Validation Loss      : 0.3832
Validation Accuracy : 0.9227
79/79 ————— 5s 51ms/step
```

	precision	recall	f1-score	support
Hate Speech	0.92	0.94	0.93	833
Offensive	0.92	0.87	0.89	833
Neither	0.93	0.95	0.94	832
accuracy			0.92	2498
macro avg	0.92	0.92	0.92	2498
weighted avg	0.92	0.92	0.92	2498

Logistic Regression Model :


```
Logistic Regression Accuracy: 0.8949
```

	precision	recall	f1-score	support
Hate Speech	0.65	0.17	0.27	286
Offensive	0.91	0.97	0.94	3838
Neither	0.85	0.82	0.83	833
accuracy			0.89	4957
macro avg	0.80	0.65	0.68	4957
weighted avg	0.88	0.89	0.88	4957

XGBoost Model :

XGBoost Accuracy: 0.8951				
	precision	recall	f1-score	support
Hate Speech	0.56	0.26	0.36	286
Offensive	0.93	0.94	0.94	3838
Neither	0.79	0.91	0.84	833
accuracy			0.90	4957
macro avg	0.76	0.70	0.71	4957
weighted avg	0.89	0.90	0.89	4957
XGBoost model and vectorizer saved.				

Combined :



Key Observations

Transformer-based models outperform traditional and sequential architectures.

Among the models implemented, BiLSTM consistently achieved higher accuracy, lower validation loss, and higher F1 scores than Logistic Regression and XGBoost. This confirms that sequential deep learning architectures with bidirectional context provide superior understanding of hate speech patterns compared to traditional machine learning baselines.

Preprocessing quality has a major impact on model performance.

Steps such as noise removal, normalization, tokenization, lemmatization, and handling class imbalance significantly improved the reliability and consistency of predictions

Despite these challenges, the system consistently demonstrated the practicality of AI-based detection in real-world conditions.

5.3 Strengths of the System

- Real-time inference capability
- Strong Contextual Understanding
- Robust Evaluation Using Multiple Metrics
- Hybrid Model Enhancement
- Adaptability for future enhancements

5.4 Limitations

Although the system performed well overall, certain limitations were observed:

1. Vulnerability to slang, sarcasm, and code-mixed language that the model was not trained on.
2. Vulnerability to Adversarial Inputs
3. Dependence on High-Quality Labelled Data.
4. Dataset imbalance can affect class-level performance.
5. Limited Generalization Across Platforms.

5.5 Discussion

The results demonstrate that deep learning-based text classification is an effective approach to identifying hate speech. BiLSTM proved suitable for real-time processing due to its speed-accuracy trade-off. The model performed consistently well across varied scenarios, validating the robustness of the dataset and training process.

The findings suggest that further improvements can be achieved through:

- Improvement of Dataset Quality and Diversity
- Class balancing
- More Advanced Text Preprocessing and Normalization
- Error Analysis and Model Refinement
- Architectural Enhancements to Existing Models

CONCLUSION

This project successfully implemented an automated Hate Speech Detection system. By comparing traditional Machine Learning models with Deep Learning, we concluded that the **Bi-Directional LSTM** offers the highest validation accuracy (~**91.6%**).

The system effectively tackles the problem of class imbalance through resampling and utilizes efficient NLP preprocessing to clean noisy social media data. This tool can serve as a backend for social media moderation bots to automatically flag or delete harmful content.

Through comprehensive experimentation involving BiLSTM, Logistic Regression, and XGBoost, it was observed that the Bi-Directional LSTM consistently outperforms traditional machine-learning approaches. The superior performance of BiLSTM highlights the critical role of sequential context in capturing nuanced linguistic patterns, implicit aggression, and semantic dependencies that are often overlooked by feature-engineering-based frameworks.

Furthermore, the incorporation of rigorous preprocessing techniques, balanced evaluation metrics, and visual performance analyses (including accuracy curves and confusion matrices) demonstrated the robustness and reliability of the proposed models. The findings also underscore the continuing challenges posed by dataset imbalance, cultural variability in hate expressions, and the need for

interpretability in automated content-moderation systems.

In conclusion, this project affirms that deep learning provides a highly promising foundation for the development of scalable, accurate, and context-aware hate speech detection systems. With further advancements in dataset diversity, model transparency, and real-time deployment methodologies, such systems have the potential to significantly enhance the safety and integrity of digital communication environments.

FUTURE WORK

Although the proposed deep learning framework demonstrates strong performance in hate speech detection, several avenues for future research remain open:

1. **Expansion to Multilingual and Code-Mixed Data:**
Real-world social media platforms contain diverse linguistic expressions, including Hinglish and other code-mixed forms. Extending the model to support multilingual datasets will enhance generalizability across regions and cultures.
2. **Handling Implicit, Sarcastic, and Context-Dependent Hate Speech:**
Future models can incorporate pragmatic understanding, sentiment shifts, and conversational context to better detect subtle or implicit hate expressions that are not explicitly abusive.
3. **Integration of Multimodal Inputs:**
Many online posts include images, memes, videos, and audio. Combining textual analysis with computer vision and speech processing models can significantly improve the detection of multimodal hate content.
4. **Bias Mitigation and Fairness Optimization:**
Future studies should examine potential demographic or linguistic biases within datasets and apply fairness-aware training methods to ensure equitable model behavior across different user groups.
5. **Deployment of Real-Time Moderation Systems:**
Integrating the model into scalable APIs or streaming pipelines will allow platforms to perform live filtering of abusive content. Further research is needed to ensure low-latency inference and system reliability.

6. Continuous Learning and Adaptive Models:

Language on social platforms evolves rapidly. Future models can incorporate online learning or periodic retraining strategies to adapt to emerging hate terms, trends, and coded vocabulary.

7. Larger and More Balanced Annotated Datasets:

Collecting high-quality, balanced datasets with expert annotations will reduce class imbalance issues and improve classification performance across minority hate categories.

REFERENCES

Davidson, T., et al. (2017). "Automated Hate Speech Detection and the Problem of Offensive Language." ICWSM.

TensorFlow Documentation. (2023). Recurrent Neural Networks (RNN) with Keras. <https://www.tensorflow.org/>

Scikit-learn Developers. (2023). Logistic Regression and TF-IDF. <https://scikit-learn.org/>

XGBoost Documentation. (2023). XGBoost Python Package. <https://xgboost.readthedocs.io/>

Dataset Source:

Davidson, T., et al. (2017). Hate Speech and Offensive Language Dataset. GitHub. <https://github.com/t-davidson/hate-speech-and-offensive-language>